SQL: The Query Language

R & G - Chapter 5



Review



- Relational Algebra (Operational Semantics)
 - Compose "tree" of operators to answer query
 - Used for query plans
- Relational Calculus (Declarative Semantics)
 - Describe what a query's answer set shall include
- Simple and powerful models for query languages

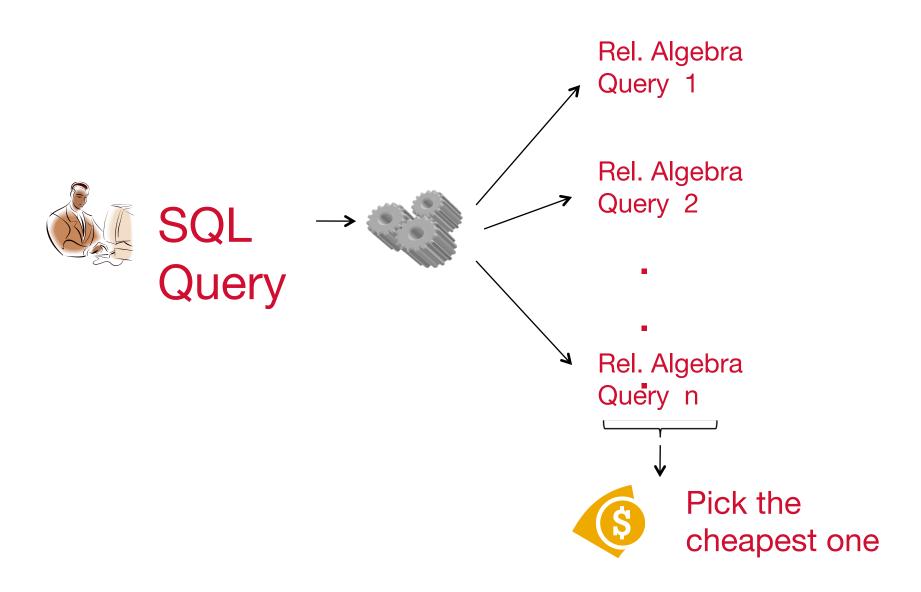
Expressivity



- An important question:
 - Just what is "sayable" in a query language?
 - This is a question of computational complexity!
- Codd's Theorem
 - relational algebra and relational calculus have equivalent expressive power
 - i.e. we can compile a declarative calculus query into an operational algebra query
- SQL adds more power
 - can be captured in variants of algebra
 - 1 key difference: multisets rather than sets
 - i.e. #duplicates in a table carefully accounted for!

Relational Query Languages





Relational Query Languages



- Two sublanguages:
 - DDL Data Definition Language
 - Define and modify schema
 - DML Data Manipulation Language
 - Queries can be written intuitively.
- DBMS is responsible for efficient evaluation.
 - The key: precise semantics for relational queries,
 Codd's Theorem
 - Optimizer can re-order operations
 - Won't affect query answer.

The SQL Query Language



- The most widely used relational query language.
- Standardized
 - (although most systems add their own "special sauce" -- including PostgreSQL)
- We will study basic constructs

Example Database



Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Boats

bid	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

Reserves

sid	bid	day
1	102	9/12/2015
2	102	9/13/2015

The SQL DDL



```
CREATE TABLE Sailors (
   sid INTEGER,
   sname CHAR(20),
   rating INTEGER,
   age REAL,
   PRIMARY KEY (sid));
CREATE TABLE Boats (
   bid INTEGER,
   bname CHAR (20),
   color CHAR(10),
   PRIMARY KEY (bid));
 CREATE TABLE Reserves (
   sid INTEGER,
   bid INTEGER,
   day DATE,
  PRIMARY KEY (sid, bid, day),
  FOREIGN KEY (sid) REFERENCES Sailors,
  FOREIGN KEY (bid) REFERENCES Boats);
```

<u>sid</u>	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

√			
<u>bid</u>	bname		color
101	Nina		red
102	Pinta		blue
103	Santa	Maria	red

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13

The SQL DML Sailors



sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Find all 18-year-old sailors:

To find just names and ratings, replace the first line:

SELECT S.sname, S.rating

Querying Multiple Relations **Berkeley**



SELECT S.sname

FROM Sailors AS S, Reserves AS R

WHERE S.sid=R.sid AND R.bid=102

Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Reserves

sid	bid	day
1	102	9/12
2	102	9/13

Basic SQL Query



<u>DISTINCT</u>: optional. Answer should not contain duplicates.

SQL default: duplicates are <u>not</u> eliminated! (Result a "multiset")

<u>target-list</u>: List of expressions over attributes of tables in *relation-list*

SELECT [DISTINCT] target-list FROM relation-list WHERE qualification

qualification: Comparisons combined using AND, OR and NOT. Comparisons are Attr *op* const or Attr1 *op* Attr2, where *op* is one of $=,<,>,\neq$, etc.

<u>relation-list</u>: List of relation names, possibly with a <u>range-variable</u> "AS" clause after each name

Query Semantics

SELECT [DISTINCT] target-list

FROM relation-list

WHERE qualification

- 1. FROM: compute cross product of tables.
- 2. WHERE: Check conditions, discard tuples that fail.
- 3. SELECT: Delete unwanted fields.
- 4. DISTINCT (optional): eliminate duplicate rows.
- Note: likely a terribly inefficient strategy!
 - Query optimizer will find more efficient plans.

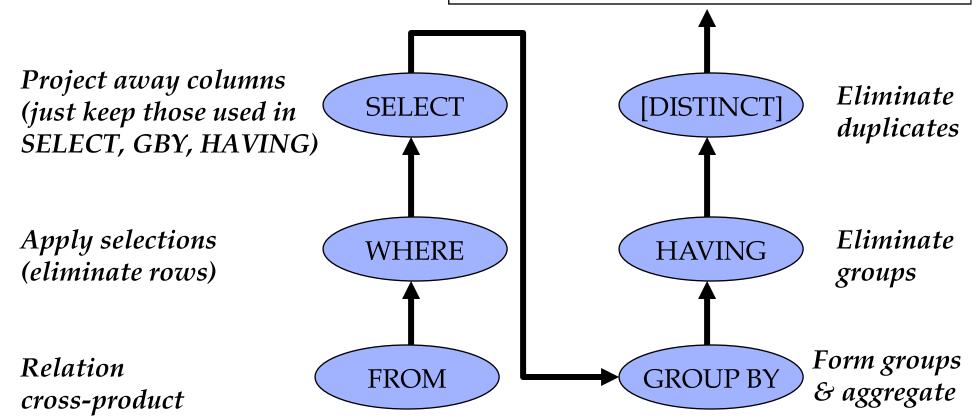
Conceptual SQL Evaluation



SELECT [DISTINCT] target-list FROM relation-list WHERE qualification

GROUP BY grouping-list

HAVING group-qualification



Find sailors who've reserved at least one boat

```
SELECT S.sid
FROM Sailors AS S, Reserves AS R
WHERE S.sid=R.sid
```

Would DISTINCT make a difference here?

About Range Variables



- Needed when ambiguity could arise.
 - e.g., same table used multiple times in FROM ("self-join")

```
SELECT x.sname, x.age, y.sname, y.age
FROM Sailors AS x, Sailors AS y
WHERE x.age > y.age
```

Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Arithmetic Expressions



```
SELECT S.age, S.age-5 AS age1, 2*S.age AS age2
FROM Sailors AS S
WHERE S.sname = 'dustin'
```

```
SELECT S1.sname AS name1, S2.sname AS name2
FROM Sailors AS S1, Sailors AS S2
WHERE 2*S1.rating = S2.rating - 1
```

String Comparisons



```
SELECT S.sname
FROM Sailors S
WHERE S.sname LIKE 'B_%B'
```

'_' stands for any one character and '%' stands for 0 or more arbitrary characters.

Most DBMSs now support standard regex as well (incl. PostgreSQL)

Find sid's of sailors who've reserved a red or a green boat



```
SELECT R.sid

FROM Boats B, Reserves R

WHERE R.bid=B.bid AND

(B.color='red' OR

B.color='green')
```

... or:

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND
B.color='red'
UNION
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='green'
```

Find sid's of sailors who've reserved a red and a green boat



```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND
(B.coTor='red' AND B.coTor='green')
```

Find sid's of sailors who've reserved a red and a green boat



```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid
        AND R.bid=B.bid
        AND B.color='red'
INTERSECT
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid
        AND R.bid=B.bid
         AND B.color='green'
```

Find sid's of sailors who've reserved a red and a green boat



Could use a self-join:

Find sid's of sailors who have not reserved a boat



```
SELECT S.sid FROM Sailors S
```

EXCEPT

```
SELECT S.sid
FROM Sailors S, Reserves R
WHERE S.sid=R.sid
```

Nested Queries: IN



Names of sailors who've reserved boat #102:

```
SELECT S.sname
FROM Sailors S
WHERE S.sid IN
(SELECT R.sid
FROM Reserves R
WHERE R.bid=102)
```

Nested Queries: NOT IN



Names of sailors who've **not** reserved boat #103:

```
SELECT S.sname
FROM Sailors S
WHERE S.sid NOT IN
(SELECT R.sid
FROM Reserves R
WHERE R.bid=103)
```

Nested Queries with Correlation



Names of sailors who've reserved boat #102:

```
SELECT S.sname

FROM Sailors S

WHERE EXISTS

(SELECT *

FROM Reserves R

WHERE R.bid=102 AND S.sid=R.sid)
```

- Subquery must be recomputed for each Sailors tuple.
 - Think of subquery as a function call that runs a query

More on Set-Comparison Operators



- we've seen: IN, EXISTS
- can also have: NOT IN, NOT EXISTS
- other forms: op ANY, op ALL

Find sailors whose rating is greater than that of some sailor called Fred:

```
SELECT *
FROM Sailors S
WHERE S.rating > ANY
   (SELECT S2.rating
   FROM Sailors S2
   WHERE S2.sname='Fred')
```

A Tough One



Find sailors who've reserved all boats.

SELECT S.sname Sailors S such that ...

FROM Sailors S

WHERE NOT EXISTS (SELECT B.bid there is no boat B

FROM Boats B without ...

WHERE NOT EXISTS (SELECT R.bid

FROM Reserves R

a Reserves tuple showing S reserved B WHERE R.bid=B.bid
AND R.sid=S.sid))

ARGMAX?



- The sailor with the highest rating
 - what about ties for highest?!

```
SELECT *
FROM Sailors S
WHERE S.rating >= ALL
   (SELECT S2.rating
   FROM Sailors S2)
```

```
SELECT *
FROM Sailors S
WHERE S.rating =
(SELECT MAX(S2.rating)
    FROM Sailors S2)
```

```
SELECT *
FROM Sailors S
ORDER BY rating DESC
LIMIT 1;
```

Null Values



- Field values are sometimes unknown or inapplicable
 - SQL provides a special value null for such situations.
- The presence of null complicates many issues. E.g.:
 - Special syntax "IS NULL" and "IS NOT NULL"
 - Assume rating = NULL. Consider predicate "rating>8".
 - True? False?
 - What about AND, OR and NOT connectives?
 - SUM?
 - We need a 3-valued logic (true, false and unknown).
 - Meaning of constructs must be defined carefully. (e.g., WHERE clause eliminates rows that don't evaluate to true.)
 - New operators (in particular, outer joins) possible/needed.

Joins



```
SELECT (column_list)
FROM table_name

[INNER | {LEFT | RIGHT | FULL } {OUTER}] JOIN table_name

ON qualification_list

WHERE ...
```

INNER is default

Inner/Natural Joins



SELECT s.sid, s.sname, r.bid FROM Sailors s, Reserves r WHERE s.sid = r.sid

SELECT s.sid, s.sname, r.bid FROM Sailors s **INNER JOIN** Reserves r **ON** s.sid = r.sid

all 3 are equivalent!

SELECT s.sid, s.sname, r.bid FROM Sailors s **NATURAL JOIN** Reserves r

 "NATURAL" means equi-join for each pair of attributes with the same name



SELECT s.sid, s.sname, r.bid FROM Sailors2 s INNER JOIN Reserves2 r ON s.sid = r.sid;

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

sid	<u>bid</u>	day
22	101	10/10/96
95	103	11/12/96

s.sid	s.name	r.bid	
22	Dustin		101
95	Bob		103

Left Outer Join



Returns all matched rows, <u>plus</u> <u>all unmatched rows from</u> <u>the table on the left</u> of the join clause (use nulls in fields of non-matching tuples)

SELECT s.sid, s.sname, r.bid FROM Sailors2 s LEFT OUTER JOIN Reserves2 r ON s.sid = r.sid;

Returns all sailors & bid for boat in any of their reservations

Note: no match for s.sid? r.bid IS NULL!



SELECT s.sid, s.name, r.bid FROM Sailors2 s LEFT OUTER JOIN Reserves2 r ON s.sid = r.sid

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

sid	bid	day
22	101	10/10/96
95	103	11/12/96

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103
31	Lubber	

Right Outer Join



 Right Outer Join returns all matched rows, <u>plus all</u> <u>unmatched rows from the table on the right</u> of the join clause

SELECT r.sid, b.bid, b.bname FROM Reserves2 r RIGHT OUTER JOIN Boats2 b ON r.bid = b.bid;

- Returns all boats & information on which ones are reserved.
- No match for b.bid? r.sid IS NULL!



SELECT r.sid, b.bid, b.bname FROM Reserves2 r RIGHT OUTER JOIN Boats2 b ON r.bid = b.bid;

sid	<u>bid</u>	day
22	101	10/10/96
95	103	11/12/96

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

r.sid	b.bid		b.name
2:	2	101	Interlake
		102	Interlake
9:	5	103	Clipper
		104	Marine

Full Outer Join



 Full Outer Join returns all (matched or unmatched) rows from the tables on both sides of the join clause

SELECT r.sid, b.bid, b.bname FROM Reserves2 r FULL OUTER JOIN Boats2 b ON r.bid = b.bid

- Returns all boats & all information on reservations
- No match for r.bid?
 - b.bid IS NULL AND b.bname IS NULL!
- No match for b.bid?
 - r.sid IS NULL!

SELECT r.sid, b.bid, b.name FROM Reserves2 r FULL OUTER JOIN Boats2 b



ON r.bid = b.bid

sid	<u>bid</u>	day
22	101	10/10/96
95	103	11/12/96

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine

Note: in this case it is the same as the ROJ! bid is a foreign key in reserves, so all reservations must have a corresponding tuple in boats.

Views: Named Queries



CREATE VIEW view_name AS select_statement

Makes development simpler
Often used for security
Not "materialized"

CREATE VIEW Reds
AS SELECT B.bid, COUNT (*) AS scount
FROM Boats2 B, Reserves2 R
WHERE R.bid=B.bid AND B.color='red'
GROUP BY B.bid

Views Instead of Relations in Queries Berkeley



CREATE VIEW Redcount AS SELECT B.bid, COUNT (*) AS scount FROM Boats2 B, Reserves2 R WHERE R.bid=B.bid AND B.color='red' GROUP BY B.bid

bid		scount		D 1
	102		1	Reds

SELECT bname, scount FROM **Redcount R**, Boats2 B WHERE R.bid=B.bid AND scount < 10

Subqueries in FROM



```
SELECT bname, scount
FROM Boats2 B,
(SELECT B.bid, COUNT (*)
FROM Boats2 B, Reserves2 R
WHERE R.bid=B.bid AND B.color='red'
GROUP BY B.bid) AS Reds(bid, scount)
WHERE Reds.bid=B.bid
AND scount < 10
```

WITH (common table expression)

```
Berkeley cs186
```

```
WITH Reds(bid, scount) AS
(SELECT B.bid, COUNT (*)
FROM Boats2 B, Reserves2 R
WHERE R.bid=B.bid AND B.color='red'
GROUP BY B.bid)
SELECT bname, scount
FROM Boats2 B, Reds
WHERE Reds.bid=B.bid
AND scount < 10
```

Discretionary Access Control



GRANT privileges ON object TO users [WITH GRANT OPTION]

- Object can be a Table or a View
- Privileges can be:
 - Select
 - Insert
 - Delete
 - References (cols) allow to create a foreign key that references the specified column(s)
 - All
- Can later be REVOKEd
- Users can be single users or groups
- See Chapter 17 for more details.

Two more important topics



Constraints

SQL embedded in other languages

Integrity Constraints



- IC conditions that every <u>legal</u> instance of a relation must satisfy.
 - Inserts/deletes/updates that violate ICs are disallowed.
 - Can ensure application semantics (e.g., sid is a key),
 - ...or prevent inconsistencies (e.g., sname has to be a string, age must be < 200)
- Types of IC's: Domain constraints, primary key constraints, foreign key constraints, general constraints.
 - Domain constraints: Field values must be of right type.
 Always enforced.
 - Primary key and foreign key constraints: coming right up.

Where do ICs Come From?




- Semantics of the real world!
- Note:
 - We can check IC violation in a DB instance
 - We can NEVER infer that an IC is true by looking at an instance.
 - An IC is a statement about all possible instances!
 - From example, we know name is not a key, but the assertion that sid is a key is given to us.
- Key and foreign key ICs are the most common
- More general ICs supported too.

Keys



- Keys are a way to associate tuples in different relations
- Keys are one form of IC



Primary Keys



- A set of fields is a superkey if:
 - No two distinct tuples can have same values in all key fields
- A set of fields is a key for a relation if it is minimal:
 - It is a superkey
 - No subset of the fields is a superkey
- what if >1 key for a relation?
 - One of the keys is chosen (by DBA) to be the primary key.
 Other keys are called candidate keys.
- E.g.
 - sid is a key for Students.
 - What about name?
 - The set {sid, gpa} is a superkey.

Primary and Candidate Keys



- Possibly many <u>candidate keys</u> (specified using UNIQUE), one of which is chosen as the <u>primary key</u>.
- Keys must be used carefully!

```
CREATE TABLE Enrolled1 CREATE TABLE Enrolled2 (sid CHAR(20), cid CHAR(20), cid CHAR(20), grade CHAR(2), PRIMARY KEY (sid,cid))

PRIMARY KEY (sid,cid))

CREATE TABLE Enrolled2 (sid CHAR(20), cid CHAR(20), grade CHAR(20), grade CHAR(20), UNIQUE (cid, grade))
```

[&]quot;For a given student and course, there is a single grade."

Primary and Candidate Keys



```
CREATE TABLE Enrolled1
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid))

INSERT INTO enrolled1 VALUES ('1234', 'cs186', 'A+');
INSERT INTO enrolled1 VALUES ('1234', 'cs186', 'F');
INSERT INTO enrolled1 VALUES ('1234', 'cs61C', 'A+');
```

INSERT INTO enrolled2 VALUES ('1234', 'cs186', 'A+');

INSERT INTO enrolled2 VALUES ('1234', 'cs186', 'F');

INSERT INTO enrolled2 VALUES ('1234', 'cs61C', 'A+');

INSERT INTO enrolled2 VALUES ('4567', 'cs186', 'A+');

CREATE TABLE Enrolled?

"For a given student and course, there is a single grade."

Primary and Candidate Keys



```
CREATE TABLE Enrolled2
CREATE TABLE Enrolled1
                                  (sid CHAR(20),
  (sid CHAR(20),
                                   cid CHAR(20),
   cid CHAR(20),
                            VS. grade CHAR(2),
   grade CHAR(2),
                                   PRIMARY KEY (sid),
   PRIMARY KEY (sid,cid));
                                   UNIQUE (cid, grade));
  INSERT INTO enrolled1 VALUES ('1234', 'cs186', 'A+');
  INSERT INTO enrolled1 VALUES ('1234', 'cs186', 'F');
  INSERT INTO enrolled1 VALUES ('1234', 'cs61C', 'A+');
  INSERT INTO enrolled2 VALUES ('1234', 'cs186', 'A+');
  INSERT INTO enrolled2 VALUES ('1234', 'cs186', 'F');
  INSERT INTO enrolled2 VALUES ('1234', 'cs61C', 'A+');
  INSERT INTO enrolled2 VALUES ('4567', 'cs186', 'A+');
```

"Students can take only one course, and no two students in a course receive the same grade."

Foreign Keys, Referential Integrity



- Foreign key: a "logical pointer"
 - Set of fields in a tuple in one relation that `refer' to a tuple in another relation.
 - Reference to primary key of the other relation.
- All foreign key constraints enforced?
 - <u>referential integrity</u>!
 - i.e., no dangling references.

Foreign Keys in SQL



- E.g. Only students listed in the Students relation should be allowed to enroll for courses.
 - sid is a foreign key referring to Students:

```
CREATE TABLE Enrolled (sid CHAR(20),cid CHAR(20),grade CHAR(2), PRIMARY KEY (sid,cid), FOREIGN KEY (sid) REFERENCES Students);
```

Enrolled

sid	cid	grade	
53666	Carnatic101	C	_
53666	Reggae203	В -	
53650	Topology112	A -	
53666	History105	B -	
11111			L
11111		A	

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Enforcing Referential Integrity

- *sid* in Enrolled: foreign key referencing Students.
- Scenarios:
 - Insert Enrolled tuple with non-existent student id?
 - Delete a Students tuple?
 - Also delete Enrolled tuples that refer to it? (CASCADE)
 - Disallow if referred to? (NO ACTION)
 - Set sid in referring Enrolled tups to a default value? (SET DEFAULT)
 - Set sid in referring Enrolled tuples to null, denoting `unknown' or `inapplicable'. (SET NULL)
- Similar issues arise if primary key of Students tuple is updated.

General Constraints

(sid INTEGER, Berkeley sname CHAR(10), rating INTEGER, age REAL, PRIMARY KEY (sid), CHECK (rating >= 1 AND rating <= 10))

- Useful when more general ICs than keys are involved.
- Can use queries to express constraint.
- Checked on insert or update.
- Constraints can be named.

```
CREATE TABLE Reserves
( sname CHAR(10),
bid INTEGER,
day DATE,
PRIMARY KEY (bid,day),
CONSTRAINT noInterlakeRes
CHECK ('Interlake' <>
( SELECT B.bname
FROM Boats B
WHERE B.bid=bid)))
```

Constraints Over Multiple Relations



CREATE TABLE Sailors

(sid INTEGER,
 sname CHAR(10),
 rating INTEGER,
 age REAL,
 PRIMARY KEY (sid),
 CHECK
 ((SELECT COUNT (S.sid) FROM Sailors S)
 + (SELECT COUNT (B.bid) FROM
 Boats B) < 100)

Constraints Over Multiple Relations



CREATE TABLE Sailors

CHECK

(sid INTEGER, sname CHAR(10), rating INTEGER, age REAL, PRIMARY KEY (sid),

Number of boats plus number of sailors is < 100

- Awkward and wrong!
 - Only checks sailors!
- ASSERTION is the right solution; not associated with either table.
 - Unfortunately, not supported in many DBMS.
 - Triggers are another solution.

((SELECT COUNT (S.sid) FROM Sailors S) + (SELECT COUNT (B.bid) FROM Boats B) < 100)

CREATE ASSERTION smallClub
CHECK
((SELECT COUNT (S.sid) FROM Sailors S)
+ (SELECT COUNT (B.bid)
FROM Boats B) < 100)

Two more important topics



Constraints

SQL embedded in other languages

Writing Applications with SQL Berkeley

- SQL is not a general purpose programming language.
 - + Tailored for data retrieval and manipulation
 - + Relatively easy to optimize and parallelize
 - Can't write entire apps in SQL alone

Options:

- Make the query language "Turing complete"
 - · Avoids the "impedance mismatch"
 - makes "simple" relational language complex
- Allow SQL to be embedded in regular programming languages.
- Q: What needs to be solved to make the latter approach work?



Cursors

- Can declare a cursor on a relation or query
- Can *open* a cursor
- Can repeatedly fetch a tuple (moving the cursor)
- Special return value when all tuples have been retrieved.
- ORDER BY allows control over the order tuples are returned.
 - Fields in ORDER BY clause must also appear in SELECT clause.
- LIMIT controls the number of rows returned (good fit w/ORDER BY)
- Can also modify/delete tuple pointed to by a cursor
 - A "non-relational" way to get a handle to a particular tuple



Database APIs

- A library with database calls (API)
 - special objects/methods
 - passes SQL strings from language, presents result sets in a language-friendly way
 - ODBC a C/C++ standard started on Windows
 - JDBC a Java equivalent
 - Most scripting languages have similar things
 - E.g. in Ruby there's the "pg" gem for Postgres
- ODBC/JDBC try to be DBMS-neutral
 - at least try to hide distinctions across different DBMSs
- Object-Relational Mappings (ORMs)
 - Ruby on Rails, Django, Spring, BackboneORM, etc.
 - Automagically map database rows into PL objects
 - Magic can be great; magic can bite you.
 - This year we won't cover ORMs much see CS169.

Summary



- Relational model has well-defined query semantics
- SQL provides functionality close to basic relational model

(some differences in duplicate handling, null values, set operators, ...)

- Typically, many ways to write a query
 - DBMS figures out a fast way to execute a query, regardless of how it is written.

Getting Serious



- Two "fancy" queries for different applications
 - Clustering Coefficient for Social Network graphs
 - Medians for "robust" estimates of the central value

Serious SQL: Social Nets Example



```
-- An undirected friend graph. Store each link once
CREATE TABLE Friends(
    fromID integer,
    toID integer,
    since date,
    PRIMARY KEY (fromID, toID),
    FOREIGN KEY (fromID) REFERENCES Users,
    FOREIGN KEY (toID) REFERENCES Users,
    CHECK (fromID < toID));</pre>
-- Return both directions
CREATE VIEW BothFriends AS
  SELECT * FROM Friends
  UNION ALL
  SELECT F.toID AS fromID, F.fromID AS toID, F.since
  FROM Friends F;
```

6 degrees of friends



Clustering Coefficient of a Node



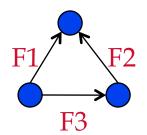
$$C_i = 2|\{e_{jk}\}|/k_i(k_i-1)$$

- where:
 - $-k_i$ is the number of neighbors of node I
 - $-e_{jk}$ is an edge between nodes j and k neighbors of i, (j < k). (A triangle!)
- I.e. Cliquishness: the fraction of your friends that are friends with each other!

 Clustering Coefficient of a graph is the average CC of all nodes.



$$C_i = 2|\{e_{jk}\}|/k_i(k_i-1)$$



CREATE VIEW NEIGHBOR_CNT AS

SELECT FROM

GROUP

CREATE VIEW TRIANGLES AS

FROM WHERE

SELECT

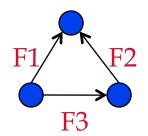
AND

AND

,



$$C_i = 2|\{e_{jk}\}|/k_i(k_i-1)$$



```
CREATE VIEW NEIGHBOR_CNT AS

SELECT fromID AS nodeID, count(*) AS friend_cnt

FROM BothFriends

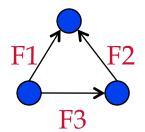
GROUP BY nodeID;
```

CREATE VIEW TRIANGLES AS

```
FROM WHERE AND AND
```

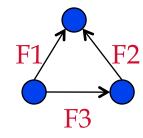


$$C_i = 2|\{e_{jk}\}| / k_i(k_i-1)$$



```
CREATE VIEW NEIGHBOR_CNT AS
SELECT fromID AS nodeID, count(*) AS friend_cnt
  FROM BothFriends
GROUP BY nodeID;
```





$$C_i = 2|\{e_{jk}\}|/k_i(k_i-1)$$

CREATE VIEW NEIGHBOR_EDGE_CNT AS

SELECT GROUP


CREATE VIEW CC_PER_NODE AS

SELECT

FROM WHERE

SELECT AVG(cc) FROM CC_PER_NODE;





$$C_i = 2|\{e_{jk}\}|/k_i(k_i-1)$$

CREATE VIEW NEIGHBOR_EDGE_CNT AS

SELECT root, COUNT(*) as cnt FROM TRIANGLES

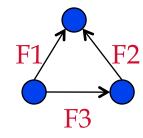
GROUP BY root;

CREATE VIEW CC_PER_NODE AS
SELECT

FROM WHERE

SELECT AVG(cc) FROM CC_PER_NODE;





$$C_i = 2|\{e_{jk}\}|/k_i(k_i-1)$$

CREATE VIEW NEIGHBOR_EDGE_CNT AS

SELECT root, COUNT(*) as cnt FROM TRIANGLES

GROUP BY root;

SELECT AVG(cc) FROM CC_PER_NODE;

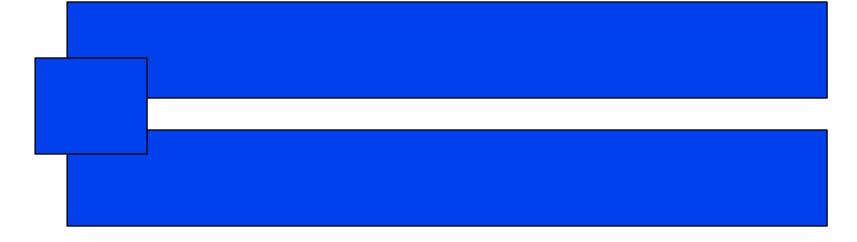
Median



- Given n values in sorted order, the one at position n/2
 - Assumes an odd # of items
 - For an even #, can take the lower of the middle 2
- A much more "robust" statistic than average
 - Q: Suppose you want the mean to be 1,000,000. What fraction of values do you have to corrupt?
 - Q2: Suppose you want the median to be 1,000,000. Same question.
 - This is called the breakdown point of a statistic.
 - Important for dealing with data outliers
 - E.g. dirty data
 - Even with real data: "overfitting"



SELECT c AS median FROM T WHERE





SELECT c AS median FROM T
WHERE
=



```
SELECT c AS median FROM T
WHERE
(SELECT COUNT(*) from T AS T1
WHERE T1.c < T.c)
=
```



```
SELECT c AS median FROM T
WHERE
(SELECT COUNT(*) from T AS T1
  WHERE T1.c < T.c)
=
(SELECT COUNT(*) from T AS T2
  WHERE T2.c > T.c);
```

Faster Median in SQL



```
SELECT x.c as median
  FROM T x, T y
 GROUP BY X.C
HAVING
 SUM(CASE WHEN y.c \ll x.c THEN 1 ELSE 0 END)
  >= (COUNT(*)+1)/2
AND
 SUM(CASE WHEN y.c >= x.c THEN 1 ELSE 0 END)
  >= (COUNT(*)/2)+1
    Why faster?
    Note: handles even # of items!
```

Using "Window Functions"



Window functions: an SQL idiom to compute with order.

http://www.postgresql.org/docs/9.3/static/tutorial-window.html

```
CREATE VIEW twocounters AS
(SELECT x,
        ROW_NUMBER() OVER (ORDER BY x ASC) AS RowAsc,
        ROW_NUMBER() OVER (ORDER BY x DESC) AS RowDesc
   FROM numbers
);
SELECT AVG(x)
FROM twocounters
WHERE RowAsc IN (RowDesc, RowDesc - 1, RowDesc + 1);
    O(n logn!)
    Note: handles even # of items.
```

Notes for Studying



- You'll be responsible for all the constructs we mentioned, except
 - Window functions
 - Programming Language APIs
- In HW3 you may write queries using:
 - Any PostgreSQL features you like
 - Except for callouts to user-defined code
 (C, Java, Python, R, etc.)