

## Entity-Relationship Diagrams and the Relational Model

R & G, Chaps. 2&3



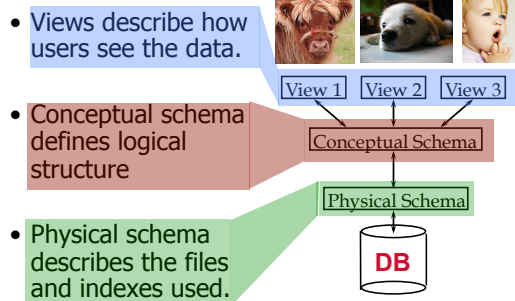
## Describing Data: Data Models

- Data model: collection of concepts for describing data.
- Schema: description of a particular collection of data, using a given data model.
- Relational model of data
  - Main concept: relation (table), rows and columns
  - Every relation has a schema
    - describes the columns
    - column names and domains



## Levels of Abstraction

Users



## Example: University Database

- Conceptual schema:
  - Students(sid text, name text, login text, age integer, gpa float)
  - Courses(cid text, cname text, credits integer)
  - Enrolled(sid text, cid text, grade text)
- Physical schema:
  - Relations stored as unordered files.
  - Index on first column of Students.
- External Schema (View):
  - Course\_info(cid text, enrollment integer)



## Data Independence


- Insulate apps from structure of data
- Logical data independence:
  - Protection from changes in *logical* structure
- Physical data independence:
  - Protection from changes in *physical* structure
- Q: Why particularly important for DBMS?

Because databases and their associated applications persist.

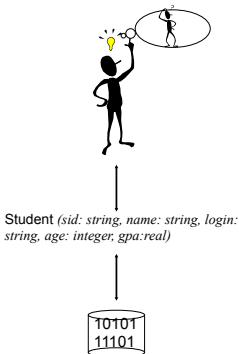



## Hellerstein's Inequality

$$\frac{dapp}{dt} \ll \frac{denv}{dt}$$


 **Data Models**

- Connect concepts to bits!
- Many models exist
- We will ground ourselves in the *Relational* model
  - clean and common
  - generalization of key/value
- *Entity-Relationship* model also handy for design
  - Translates down to Relational




 **Why Focus on Relational Model?**


- Most widely used model.
- Other models exist (and co-exist)
  - “Legacy systems” in older models
    - e.g., IBM’s IMS
  - Object-Relational mergers
    - Object-Oriented features provided by DBMS
    - Object-Relational Mapping (ORM) outside the DBMS
      - A la Rails (Ruby), Django (Python), Hibernate (Java)
  - XML, JSON, etc.
    - Nested and “semi-structured” data
    - Languages like Xquery, XSLT, JSONic
    - Most relational engines now handle these to a degree

 **Entity-Relationship Model**


- Relational model is a great formalism
  - and a clean system framework
- But a bit detailed for design time
  - a bit fussy for brainstorming
  - hard to communicate to customers
- Entity-Relationship model is a popular “shim” over relational model
  - graphical, slightly higher level

 **Steps in Traditional Database Design**

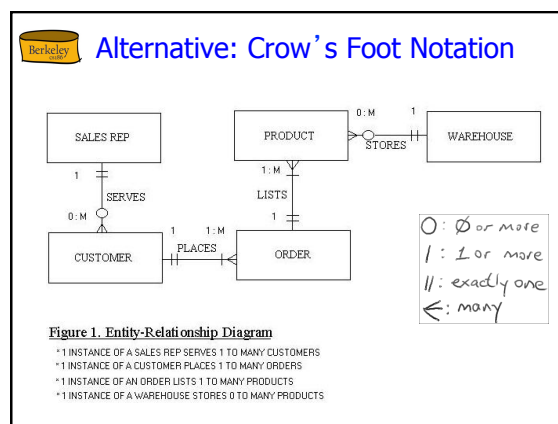
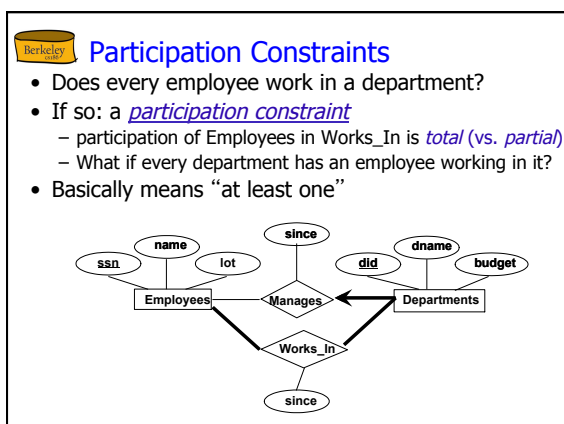
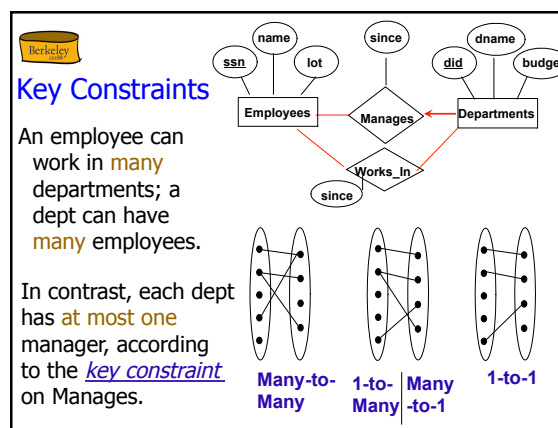
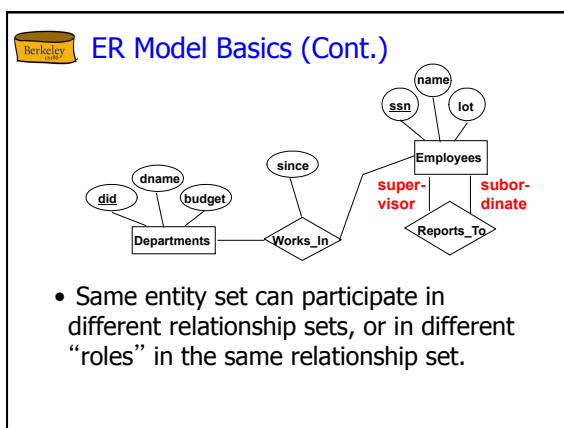
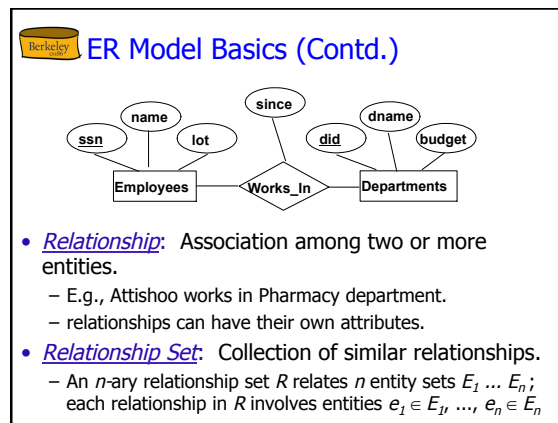
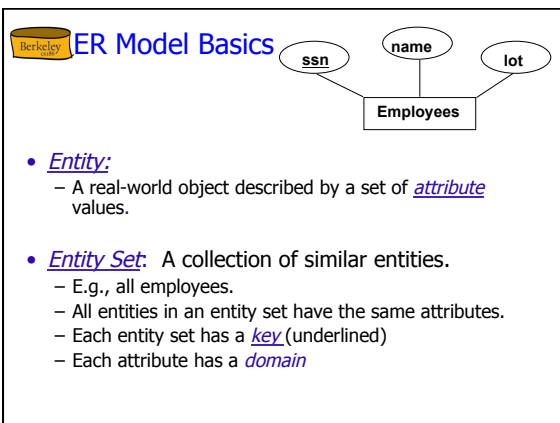
- **Requirements Analysis**
  - user needs; what must database do?
- **Conceptual Design**
  - high level description (often done w/ER model)
  - *Rails encourages you to program here*
- **Logical Design**
  - translate ER into DBMS data model
  - *Rails requires you to help here too*
- **Schema Refinement**
  - consistency, normalization
- **Physical Design** - indexes, disk layout
- **Security Design** - who accesses what, and how

 **Conceptual Design**

- What are the entities and relationships?
- What info about E’s & R’s should be in DB?
- What *integrity constraints (business rules)* hold?
- *ER diagram* is the “schema”
- Can map an ER diagram into a relational schema.
- Conceptual design is where the SW/data engineering *begins*
  - Rails “models”

 **Modern pattern: “Schema on Use”**

- What about more agile, less governed environments?
- Don’t let the lack of schema prevent storing data!
  - Just use binary, text, CSV, JSON, xlsx, etc.
  - Can shove into a DBMS, or just a filesystem (e.g. HDFS)
  - Most database engines can query files directly these days
- Wrangle the data into shape as needed
  - Essentially defining views over the raw data
  - This amounts to database design, at the view level
  - What about integrity constraints?
    - Instead, define “anomaly indicator” columns – or queries
- Fits well with read/append-only data
  - E.g. Big Data, a la Hadoop
  - Less of a fit with update-heavy data
- Analogies to strong vs. loose typing in PL





## Summary so far

- Entities and Entity Set (boxes)
- Relationships and Relationship sets (diamonds)
- Key constraints (arrows)
- Participation constraints (bold for Total)


These are enough to get started, but we'll need more...



## Weak Entities

A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.

- Owner entity set and weak entity set must participate in a one-to-many relationship set (one owner, many weak entities).
- Weak entity set must have total participation in this *identifying* relationship set.



Weak entities have only a “partial key” (dashed underline)

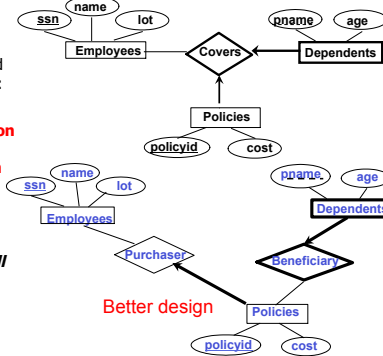


## Binary vs. Ternary Relationships

If each policy is owned by just 1 employee:

**Key constraint on Policies would mean policy can only cover 1 dependent!**

- Think through *all* the constraints in the 2nd diagram!




## Binary vs. Ternary Relationships (Contd.)

- Previous example:
  - 2 binary relationships better than 1 ternary relationship.
- An example in the other direction:
  - ternary relationship set *Contracts* relates entity sets *Parts*, *Departments* and *Suppliers*
  - relationship set has descriptive attribute *qty*.
  - no combo of binary relationships is a substitute!
    - See next slide...



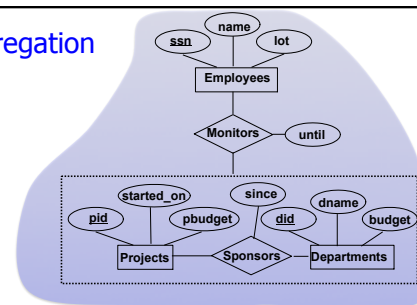
## Binary vs. Ternary Relationships (Contd.)



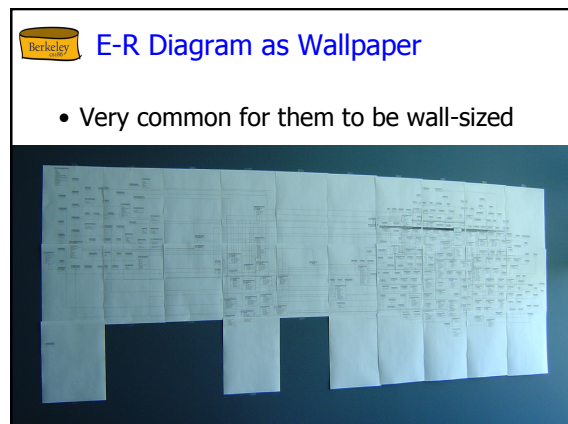
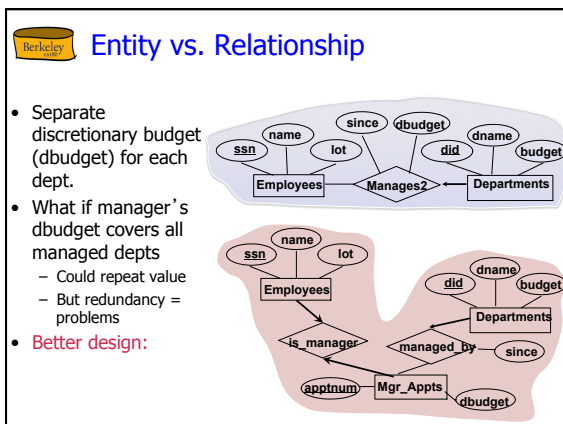
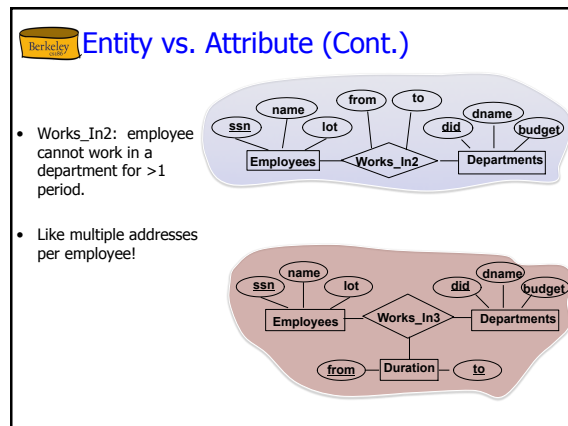
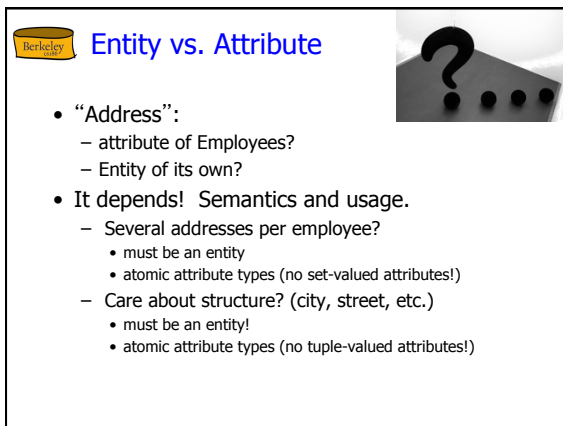
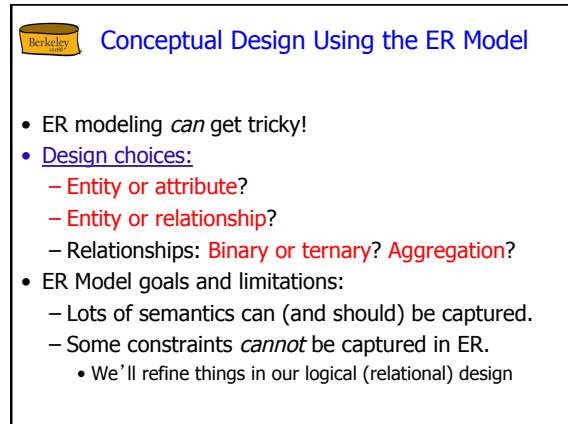

- S “can-supply” P, D “needs” P, and D “deals-with” S does not imply that D has agreed to buy P from S.
- How do we record *qty*?



## Aggregation



Allows relationships with *relationship sets*.






## Converting ER to Relational

- Fairly analogous structure
- But many simple concepts in ER are subtle to specify in relations



## Logical DB Design: ER to Relational

- Entity sets to tables.



ssn	name	lot
123-22-3666	Attishoo	48
231-31-5368	Smiley	22
131-24-3650	Smethurst	35

```
CREATE TABLE Employees
(ssn CHAR(11),
 name CHAR(20),
 lot INTEGER,
 PRIMARY KEY (ssn))
```



## Relationship Sets to Tables

- In translating a **many-to-many** relationship set to a relation, attributes of the relation must include:

- Keys for each participating entity set (as foreign keys). This set of attributes forms a **superkey** for the relation.
- All descriptive attributes.



```
CREATE TABLE Works_In(
ssn CHAR(11),
 did INTEGER,
 since DATE,
 PRIMARY KEY (ssn, did),
 FOREIGN KEY (ssn)
 REFERENCES Employees,
 FOREIGN KEY (did)
 REFERENCES Departments)
```

ssn	did	since
123-22-3666	51	1/1/91
123-22-3666	56	3/3/93
231-31-5368	51	2/2/92




## Review: Key Constraints


- Each dept has at most one manager, according to the **key constraint** on Manages.




1-to-1



1-to Many



Many-to-1




Many-to-Many

Translation to relational model?



## Translating ER with Key Constraints



- Since each department has a unique manager, we could instead combine Manages and Departments.

```
CREATE TABLE Manages(
ssn CHAR(11),
 did INTEGER,
 since DATE,
 PRIMARY KEY (did),
 FOREIGN KEY (ssn)
 REFERENCES Employees,
 FOREIGN KEY (did)
 REFERENCES Departments)
```


Vs.

```
CREATE TABLE Dept_Mgr(
did INTEGER,
 dname CHAR(20),
 budget REAL,
 ssn CHAR(11),
 since DATE,
 PRIMARY KEY (did),
 FOREIGN KEY (ssn)
 REFERENCES Employees)
```



## Review: Participation Constraints

- Does every department have a manager?
  - If so, this is a **participation constraint**: the participation of Departments in Manages is said to be **total (vs. partial)**.
    - Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!)




## Participation Constraints in SQL

- We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to CHECK constraints).

```
CREATE TABLE Dept_Mgr(
  did INTEGER,
  dname CHAR(20),
  budget REAL,
  ssn CHAR(11) NOT NULL,
  since DATE,
  PRIMARY KEY (did),
  FOREIGN KEY (ssn) REFERENCES
Employees
ON DELETE NO ACTION)
```

## Review: Weak Entities

- A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
  - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
  - Weak entity set must have total participation in this *identifying* relationship set.



## Translating Weak Entity Sets

- Weak entity set and identifying relationship set are translated into a single table.
  - When the owner entity is deleted, all owned weak entities must also be deleted.

```
CREATE TABLE Dep_Policy (
  pname CHAR(20),
  age INTEGER,
  cost REAL,
  ssn CHAR(11) NOT NULL,
  PRIMARY KEY (pname, ssn),
  FOREIGN KEY (ssn) REFERENCES Employees
ON DELETE CASCADE)
```

## Summary of Conceptual Design

- Conceptual design* follows *requirements analysis*,
  - Yields a high-level description of data to be stored
  - You may want to postpone it for read-only "schema on use"
- ER model popular for conceptual design
  - Constructs are expressive, close to the way people think about their applications.
  - Note: There are many variations on ER model
    - Both graphically and conceptually
- Basic constructs: *entities*, *relationships*, and *attributes* (of entities and relationships).
- Some additional constructs: *weak entities*, *ISA hierarchies* (see text if you're curious), and *aggregation*.

## Summary of ER (Cont.)

- Several kinds of integrity constraints:
  - key constraints*
  - participation constraints*
- Some *foreign key constraints* are also implicit in the definition of a relationship set.
- Many other constraints (notably, *functional dependencies*) cannot be expressed.
- Constraints play an important role in determining the best database design for an enterprise.

## Summary of ER (Cont.)

- ER design is *subjective*. There are often many ways to model a given scenario!
- Analyzing alternatives can be tricky, especially for a large enterprise. Common choices include:
  - Entity vs. attribute, entity vs. relationship, binary or n-ary relationship, whether or not to use ISA hierarchies, aggregation.
- Ensuring good database design: resulting relational schema should be analyzed and refined further.
  - Functional Dependency information and normalization techniques are especially useful.